

1. *The nonplanning method looks particularly poor in Figure 8.3 because it is a one-step method; a method using multi-step bootstrapping would do better. Do you think one of the multi-step bootstrapping methods from Chapter 7 could do as well as the Dyna method? Explain why or why not.*

I think it can come close. An n -step bootstrapping algorithm with big enough n (or a Monte Carlo method) would update at the end of the first episode all the action-values we encountered during the episode. I expect the policy based on these updated action-values to be quite good.

2. *Why did the Dyna agent with exploration bonus, Dyna-Q+, perform better in the first phase as well as in the second phase of the blocking and shortcut experiments?*

At first, both algorithms might find an okay policy that is suboptimal. Due to the exploration, Dyna-Q+ will realize it sooner that it's a suboptimal policy.

3. *Careful inspection of Figure 8.5 reveals that the difference between Dyna-Q+ and Dyna-Q narrowed slightly over the first part of the experiment. What is the reason for this?*

After finding the optimal path, Dyna-Q always uses that path, whereas Dyna-Q+ does some exploration from time to time.

4. *Programming. The exploration bonus described above actually changes the estimated values of states and actions. Is this necessary? Suppose the bonus $\kappa\sqrt{\tau}$ was used not in updates, but solely in action selection. That is, suppose the action selected was always that for which $Q(S_t, a) + \kappa\sqrt{\tau(S_t, a)}$ was maximal. Carry out a gridworld experiment that tests and illustrates the strengths and weaknesses of this alternate approach.*

My expectation is that in some cases it's worse to use the bonus only in the action selection. Suppose that we need a series of exploratory steps to find the shortcut. If we update the actual values, then once the exploration bonus is big enough, we will soon find the shortcut. If we only use the bonus for action selection, it might happen that we take one exploratory step but before reaching the shortcut, we get back to the old path. If this happens, we will need to wait until the first exploratory step accumulates a huge bonus again to go down on that path again.

I couldn't reproduce the cumulative reward of 400 shown in the book, my implementation of Dyna-Q+ only reached 350 and I couldn't figure out why. Anyway, here is the result.

